

ASSIGNMENT SOLUTION

The following Common Boilerplate code to create a Spark Session has to be executed before running the queries.

```
from pyspark.sql import SparkSession
import getpass

username = getpass.getuser()

spark= SparkSession. \
builder. \
config('spark.ui.port','0'). \
config("spark.sql.warehouse.dir", f"/user/{username}/warehouse"). \
enableHiveSupport(). \
master('yarn'). \
getOrCreate()
```

Note : Use Pyspark2 for executing the below queries.

Question 1

A.1

```
cust_schema = 'customer_id long,purchase_date date,product_id
integer,transaction_amount double'
```

```
transactions_df = spark.read \
.format("csv") \
.schema(cust_schema) \
.load("/public/trendytech/datasets/cust_transf.csv")
```

```
start_date = "2023-05-01"
```

```
end_date = "2023-06-30"
```

```
filtered_df = transactions_df.filter((transactions_df.purchase_date >=
start_date) & (transactions_df.purchase_date <= end_date))
```

```
revenue_df =
filtered_df.groupBy("product_id").sum("transaction_amount").withColumnRenamed("sum(transaction_amount)", "revenue")
```

```
top_products_no_cache =
revenue_df.sort("revenue",ascending=False).limit(10).show()
```

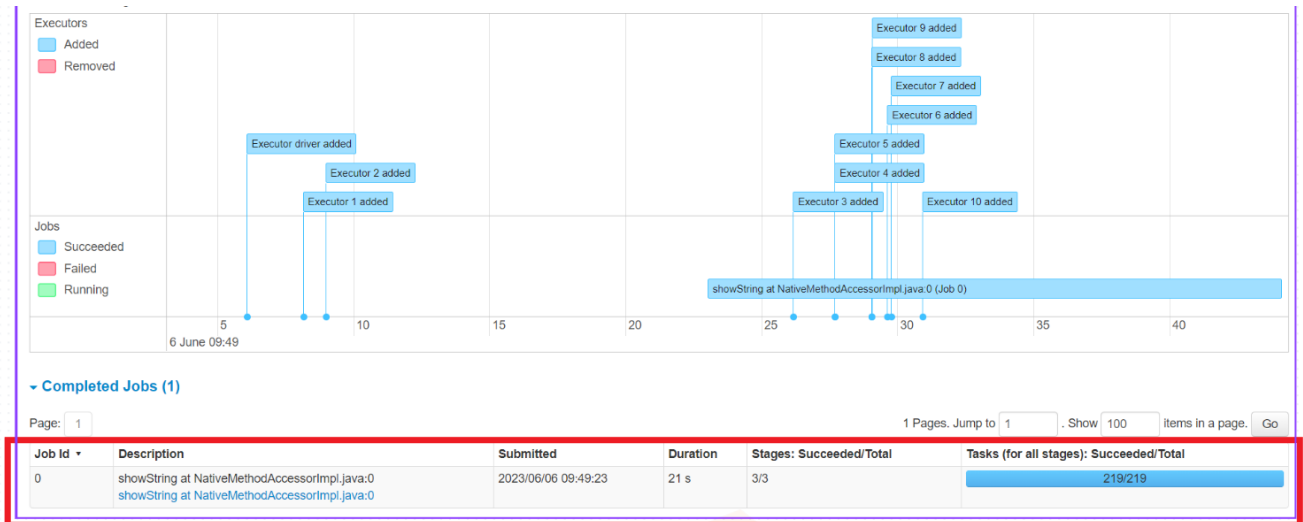
```
cust_schema = 'customer_id long,purchase_date date,product_id integer,transaction_amount double'

transactions_df = spark.read \
    .format("csv") \
    .schema(cust_schema) \
    .load("/public/trendytech/datasets/cust_transf.csv")

start_date = "2023-05-01"
end_date = "2023-06-30"
filtered_df = transactions_df.filter((transactions_df.purchase_date >= start_date) & (transactions_df.purchase_date <= end_date))

revenue_df = filtered_df.groupBy("product_id").sum("transaction_amount").withColumnRenamed("sum(transaction_amount)", "revenue")
top_products_no_cache = revenue_df.sort("revenue",ascending=False).limit(10).show()
```

product_id	revenue
1001	8.747870076028482E8
1003	6.99794607594988E8
1002	5.2480220758978057E8
1005	4.373060075933379E8
1004	3.498098075985674E8
1015	12537.909999999963
1014	11492.909999999963
1013	10447.909999999963
1012	9402.909999999965
1011	8357.909999999967



#with caching

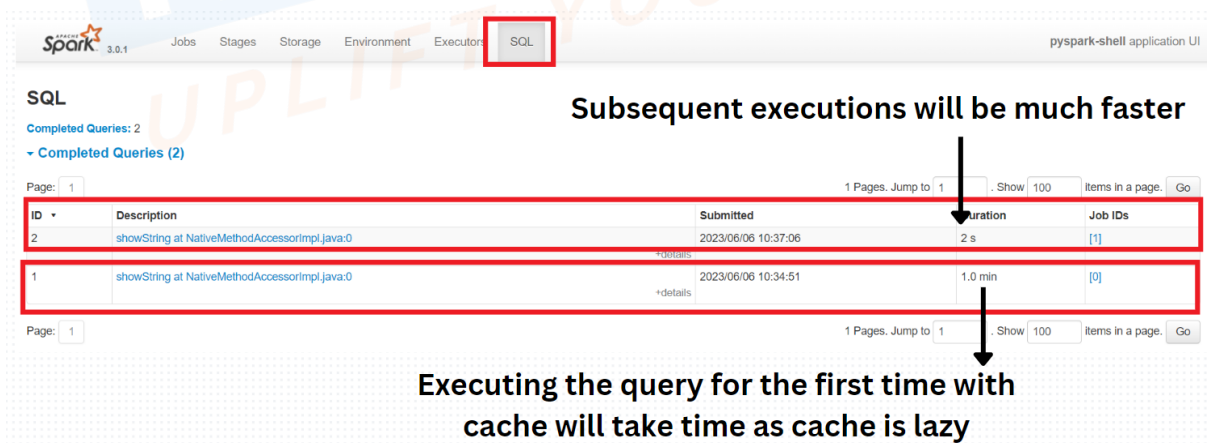
start_date = "2023-05-01"

end_date = "2023-06-30"

cached_filtered_df = transactions_df.filter((transactions_df.purchase_date >= start_date) & (transactions_df.purchase_date <= end_date)).cache()

revenue_df_with_cache =
cached_filtered_df.groupBy("product_id").sum("transaction_amount").withColumnRenamed("sum(transaction_amount)", "revenue")

top_products_with_cache = revenue_df_with_cache.orderBy("revenue", ascending=False).limit(10).show()



A.2

```
customer_transactions =  
filtered_df.groupBy("customer_id").sum("transaction_amount").withColumnRe  
named("sum(transaction_amount)", "cust_amount")
```

```
customer_transactions.show()
```

```
top_customers = customer_transactions.sort("cust_amount",  
ascending=False)
```

```
top_10_customers = top_customers.limit(10).show()
```

A.3

```
spark.sql("create database tt_cust_transaction")
```

#before caching

```
spark.sql("create table  
tt_cust_transaction.customer_transactions_ext(customer_id  
long,purchase_date date,product_id integer,transaction_amount double)  
USING csv location '/public/trendytech/datasets/cust_transf.csv'")
```

```
spark.sql("SELECT product_id, SUM(transaction_amount) AS revenue FROM  
tt_cust_transaction.customer_transactions_ext WHERE purchase_date >=  
'2023-05-01' AND purchase_date <= '2023-06-30' GROUP BY product_id  
ORDER BY revenue DESC LIMIT 10").show()
```

product_id	revenue
1001	8.747870076028483E8
1003	6.997946075949881E8
1002	5.2480220758978045E8
1005	4.373060075933379E8
1004	3.498098075985674E8
1015	12537.9099999999963
1014	11492.9099999999963
1013	10447.9099999999963
1012	9402.9099999999965
1011	8357.9099999999967

```
spark.sql("SELECT customer_id, SUM(transaction_amount) AS cust_amount
FROM tt_cust_transaction.customer_transactions_ext WHERE
purchase_date >= '2023-05-01' AND purchase_date <= '2023-06-30' GROUP
BY customer_id ORDER BY cust_amount DESC LIMIT 10").show()
```

customer_id	cust_amount
1001	3.180884580005336E8
1004	3.101342580008687E8
1005	2.6240905800151232E8
1003	2.1468385800145328E8
1002	2.0672965800144082E8
1011	1.9086143271084768E8
1006	1.9085620771084768E8
1015	1.6700301271081635E8
1010	1.6699778771081635E8
1014	1.5109356771079E8

A.4

```
from pyspark.sql.functions import year, month
```

```
from pyspark.sql.functions import countDistinct
```

```
cust_schema = 'customer_id long,purchase_date date,product_id  
integer,transaction_amount double'
```

```
transactions_df = spark.read \
```

```
.format("csv") \
```

```
.schema(cust_schema) \
```

```
.load("/public/trendytech/datasets/cust_transf.csv")
```

```
new_df = transactions_df.withColumn("purchase_year",  
year("purchase_date")).withColumn("purchase_month",  
month("purchase_date"))
```

```
customer_month_counts = new_df.groupBy("customer_id", "purchase_year",  
"purchase_month").agg(countDistinct("purchase_month").alias("distinct_months"))
```

```
regular_customers = customer_month_counts.filter("distinct_months = 1") \
```

```
.groupBy("customer_id").count() \
```

```
.orderBy("count", ascending=False).limit(10).show()
```

```
regular_customers = customer_month_counts.filter("distinct_months = 1") \
    .groupBy("customer_id").count() \
    .orderBy("count", ascending=False).limit(10).show()
```

```
+-----+-----+
|customer_id|count|
+-----+-----+
|      1009|    2|
|      1012|    2|
|      1001|    2|
|      1011|    2|
|      1007|    2|
|      1005|    2|
|      1010|    2|
|      1002|    2|
|      1006|    2|
|      1013|    2|
+-----+-----+
```

SQL

Completed Queries: 24

Completed Queries (24)

Page: 1

1 Pages. Jump to 1. Show 100 Items in a page. Go

ID	Description	Submitted	Duration	Job IDs
23	showString at NativeMethodAccessorImpl.java:0	2023/06/08 03:05:44	13 s	[7]

Time taken without cache or persist

#using persist

```
customer_month_counts = new_df.groupBy("customer_id", "purchase_year",
    "purchase_month").agg(countDistinct("purchase_month").alias("distinct_months")).persist()
```

```
regular_customers = customer_month_counts.filter("distinct_months = 1") \
    .groupBy("customer_id").count() \
    .orderBy("count", ascending=False).limit(10).show()
```


A.5

#using cache

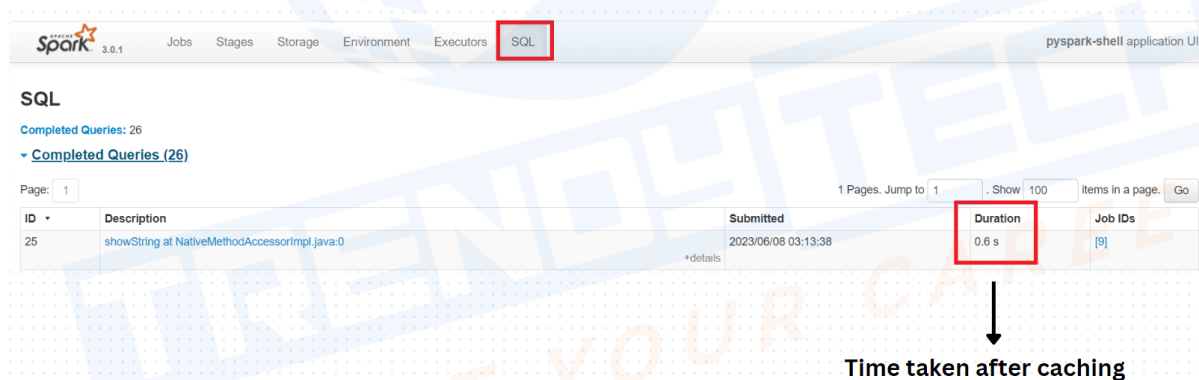
```
from pyspark.sql.functions import year, month
```

```
new_df = transactions_df.withColumn("purchase_year",  
year("purchase_date")).withColumn("purchase_month",  
month("purchase_date"))
```

```
from pyspark.sql.functions import countDistinct
```

```
customer_month_counts = new_df.groupBy("customer_id", "purchase_year",  
"purchase_month").agg(countDistinct("purchase_month").alias("distinct_months")).cache()
```

```
regular_customers = customer_month_counts.filter("distinct_months = 1") \  
    .groupBy("customer_id").count() \  
    .orderBy("count", ascending=False).limit(10).show()
```



The screenshot shows the Databricks SQL interface. At the top, there's a navigation bar with tabs for Jobs, Stages, Storage, Environment, Executors, and SQL (which is highlighted with a red box). Below the navigation bar, the 'SQL' section is active, showing 'Completed Queries: 26'. A table of completed queries is displayed, with the first query (ID 25) highlighted. The query description is 'showString at NativeMethodAccessorImpl.java:0'. The 'Submitted' time is '2023/06/08 03:13:38'. The 'Duration' is '0.6 s', which is highlighted with a red box. An arrow points from this box to the text 'Time taken after caching'.

ID	Description	Submitted	Duration	Job IDs
25	showString at NativeMethodAccessorImpl.java:0	2023/06/08 03:13:38	0.6 s	[9]

#using persist

```
from pyspark.sql.functions import year, month
```

```
from pyspark.sql.functions import countDistinct
```

```
from pyspark.storagelevel import StorageLevel
```

```
new_df = transactions_df.withColumn("purchase_year",  
year("purchase_date")).withColumn("purchase_month",  
month("purchase_date"))
```

```
customer_month_counts = new_df.groupBy("customer_id", "purchase_year",  
"purchase_month").agg(countDistinct("purchase_month").alias("distinct_months")).persist(StorageLevel.MEMORY_AND_DISK_SER)
```

```
regular_customers = customer_month_counts.filter("distinct_months = 1") \  
    .groupBy("customer_id").count() \  
    .orderBy("count", ascending=False).limit(10).show()
```

A.6

#MEMORY ONLY

```
from pyspark.sql.functions import year, month  
from pyspark.sql.functions import countDistinct  
from pyspark.storagelevel import StorageLevel
```

```
new_df = transactions_df.withColumn("purchase_year",  
year("purchase_date")).withColumn("purchase_month",  
month("purchase_date"))
```

```
customer_month_counts = new_df.groupBy("customer_id", "purchase_year",  
"purchase_month").agg(countDistinct("purchase_month").alias("distinct_months")).persist(StorageLevel.MEMORY_ONLY)
```

```
regular_customers = customer_month_counts.filter("distinct_months = 1") \  
    .groupBy("customer_id").count() \  
    .orderBy("count", ascending=False).limit(10).show()
```

APACHE spark 2.4.7 Jobs Stages Storage Environment Executors **SQL** pyspark-shell application UI

SQL

Completed Queries: 1

Completed Queries (1)

ID	Description	Submitted	Duration	Job IDs
0	showString at NativeMethodAccessorImpl.java:0	+details 2023/06/08 04:49:43	25 s	[0]

Time taken with persist - MEMORY_ONLY

MEMORY_ONLY_SER

```
from pyspark.sql.functions import year, month
from pyspark.sql.functions import countDistinct
from pyspark.storagelevel import StorageLevel
```

```
new_df = transactions_df.withColumn("purchase_year",
year("purchase_date")).withColumn("purchase_month",
month("purchase_date"))
```

```
customer_month_counts = new_df.groupBy("customer_id", "purchase_year",
"purchase_month").agg(countDistinct("purchase_month").alias("distinct_months")).persist(StorageLevel.MEMORY_ONLY_SER)
```

```
regular_customers = customer_month_counts.filter("distinct_months = 1") \
.groupBy("customer_id").count() \
.orderBy("count", ascending=False).limit(10).show()
```

APACHE spark 2.4.7 Jobs Stages Storage Environment Executors **SQL** pyspark-shell application UI

SQL

Completed Queries: 1

Completed Queries (1)

ID	Description	Submitted	Duration	Job IDs
0	showString at NativeMethodAccessorImpl.java:0	+details 2023/06/08 04:54:02	1.0 min	[0]

Time taken with persist - MEMORY_ONLY_SER

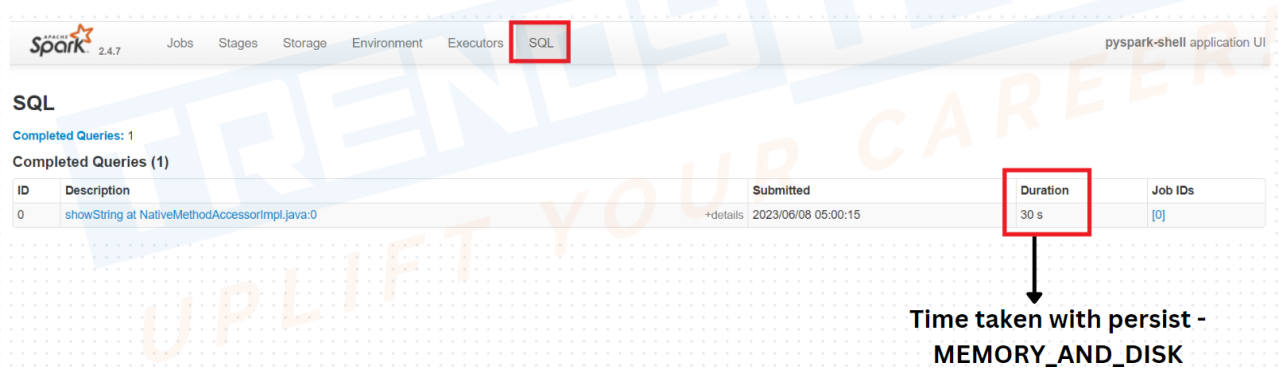
MEMORY_AND_DISK

```
from pyspark.sql.functions import year, month
from pyspark.sql.functions import countDistinct
from pyspark.storagelevel import StorageLevel
```

```
new_df = transactions_df.withColumn("purchase_year",
year("purchase_date")).withColumn("purchase_month",
month("purchase_date"))
```

```
customer_month_counts = new_df.groupBy("customer_id", "purchase_year",
"purchase_month").agg(countDistinct("purchase_month").alias("distinct_months")).persist(StorageLevel.MEMORY_AND_DISK)
```

```
regular_customers = customer_month_counts.filter("distinct_months = 1") \
    .groupBy("customer_id").count() \
    .orderBy("count", ascending=False).limit(10).show()
```



The screenshot shows the Databricks SQL interface. At the top, there's a navigation bar with tabs for Jobs, Stages, Storage, Environment, Executors, and SQL (which is highlighted with a red box). Below the navigation bar, the 'SQL' section is active. It shows 'Completed Queries: 1' and a table of completed queries. The table has columns: ID, Description, Submitted, Duration, and Job IDs. The first row shows ID 0, Description 'showString at NativeMethodAccessorImpl.java:0', Submitted '2023/06/08 05:00:15', Duration '30 s' (highlighted with a red box), and Job IDs '[0]'. An arrow points from the 'Duration' cell to the text 'Time taken with persist - MEMORY_AND_DISK'.

ID	Description	Submitted	Duration	Job IDs
0	showString at NativeMethodAccessorImpl.java:0	2023/06/08 05:00:15	30 s	[0]

Time taken with persist -
MEMORY_AND_DISK

MEMORY_AND_DISK_SER

```
from pyspark.sql.functions import year, month
from pyspark.sql.functions import countDistinct
```

```
from pyspark.storagelevel import StorageLevel
```

```
new_df = transactions_df.withColumn("purchase_year",  
year("purchase_date")).withColumn("purchase_month",  
month("purchase_date"))
```

```
customer_month_counts = new_df.groupBy("customer_id", "purchase_year",  
"purchase_month").agg(countDistinct("purchase_month").alias("distinct_months")).persist(StorageLevel.MEMORY_AND_DISK_SER)
```

```
regular_customers = customer_month_counts.filter("distinct_months = 1") \  
    .groupBy("customer_id").count() \  
    .orderBy("count", ascending=False).limit(10).show()
```

DISK_ONLY

```
from pyspark.sql.functions import year, month  
from pyspark.sql.functions import countDistinct  
from pyspark.storagelevel import StorageLevel
```

```
new_df = transactions_df.withColumn("purchase_year",  
year("purchase_date")).withColumn("purchase_month",  
month("purchase_date"))
```

```
customer_month_counts = new_df.groupBy("customer_id", "purchase_year",  
"purchase_month").agg(countDistinct("purchase_month").alias("distinct_months")).persist(StorageLevel.DISK_ONLY)
```

```
regular_customers = customer_month_counts.filter("distinct_months = 1") \  
    .groupBy("customer_id").count() \
```

```
.orderBy("count", ascending=False).limit(10).show()
```

B)

#user defined function

```
def get_customer_history(customer_id):
```

```
    customer_history_df = transactions_df.filter(transactions_df.customer_id ==  
customer_id).cache()
```

```
    return customer_history_df
```

#pass the customer_id you want to get the history

```
customer_id = 1001
```

```
customer_history_df = get_customer_history(customer_id)
```

```
customer_history_df.show()
```

C)

```
cached_filtered_df.unpersist()
```

```
spark.sql("uncache table tt_cust_transaction.customer_transactions_ext")
```

Question 2

```
spark.sql("create database tt_assignments_hotel_usecase")
```

```
spark.sql("CREATE TABLE
```

```
tt_assignments_hotel_usecase.hotel_bookings_external (booking_id INT,  
guest_name STRING, checkin_date DATE, checkout_date DATE, room_type  
STRING, total_price DOUBLE) USING csv location  
'/public/trendytech/datasets/hotel_data.csv' ")
```

```
spark.sql("select * from
```

```
tt_assignments_hotel_usecase.hotel_bookings_external limit 5").show()
```

```
spark.sql("create database tt_assignments_hotel_usecase")
```

```
spark.sql("CREATE TABLE tt_assignments_hotel_usecase.hotel_bookings_external (booking_id INT,
```

```
spark.sql("select * from tt_assignments_hotel_usecase.hotel_bookings_external limit 5").show()
```

booking_id	guest_name	checkin_date	checkout_date	room_type	total_price
1	John Doe	2023-05-01	2023-05-05	Standard	400.0
2	Jane Smith	2023-05-02	2023-05-06	Deluxe	600.0
3	Mark Johnson	2023-05-03	2023-05-08	Standard	450.0
4	Sarah Wilson	2023-05-04	2023-05-07	Executive	750.0
5	Emily Brown	2023-05-06	2023-05-09	Deluxe	550.0

A)

```
count_before_caching = spark.sql("SELECT COUNT(*) FROM  
tt_assignments_hotel_usecase.hotel_bookings_external").show()
```

```
count_before_caching = spark.sql("SELECT COUNT(*) FROM tt_assignments_hotel_usecase.hotel_bookings_external").show()
```

count(1)
107

B)

```
avg_price_without_caching = spark.sql("SELECT room_type, AVG(total_price)  
FROM tt_assignments_hotel_usecase.hotel_bookings_external GROUP BY  
room_type limit 100").show()
```

```
avg_price_without_caching = spark.sql("SELECT room_type, AVG(total_price) FROM tt_assignments
```

room_type	avg(total_price)
Executive	750.0
Deluxe	575.5813953488372
Standard	425.0

#with caching

```
spark.sql("cache table tt_hotel.hotel_bookings_external")
```

A)

```
count_after_caching = spark.sql("SELECT COUNT(*) FROM  
tt_hotel.hotel_bookings_external").show()
```

B)

```
avg_price_with_caching = spark.sql("SELECT room_type, AVG(total_price)  
FROM tt_hotel.hotel_bookings_external GROUP BY room_type limit  
100").show()
```

Completed Jobs (6)

Page: 1 1 Pages, Jump to 1 . Show 100 Items in a page. Go

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
5	showString at NativeMethodAccessorImpl.java:0 showString at NativeMethodAccessorImpl.java:0	2023/06/06 08:21:50	0.6 s	3/3	202/202
4	showString at NativeMethodAccessorImpl.java:0 showString at NativeMethodAccessorImpl.java:0	2023/06/06 08:21:18	32 ms	2/2	2/2
3	sql at NativeMethodAccessorImpl.java:0 sql at NativeMethodAccessorImpl.java:0	2023/06/06 08:20:38	0.2 s	2/2	2/2
2	showString at NativeMethodAccessorImpl.java:0 showString at NativeMethodAccessorImpl.java:0	2023/06/06 08:16:03	1 s	3/3	202/202
1	showString at NativeMethodAccessorImpl.java:0 showString at NativeMethodAccessorImpl.java:0	2023/06/06 08:13:16	2 s	2/2	2/2
0	showString at NativeMethodAccessorImpl.java:0 showString at NativeMethodAccessorImpl.java:0	2023/06/06 08:09:03	2 s	2/2	2/2

Before Caching → After Caching

Job	Without caching	With caching
1	1s	0.6s
2	0.2s	32ms

Note: You can see a large difference when dealing with really big data. Here since the data is small, the comparisons might be very less and might be varying.

C)

```
spark.sql("uncache table tt_hotel.hotel_bookings_external")
```