

Spark on YARN Architecture

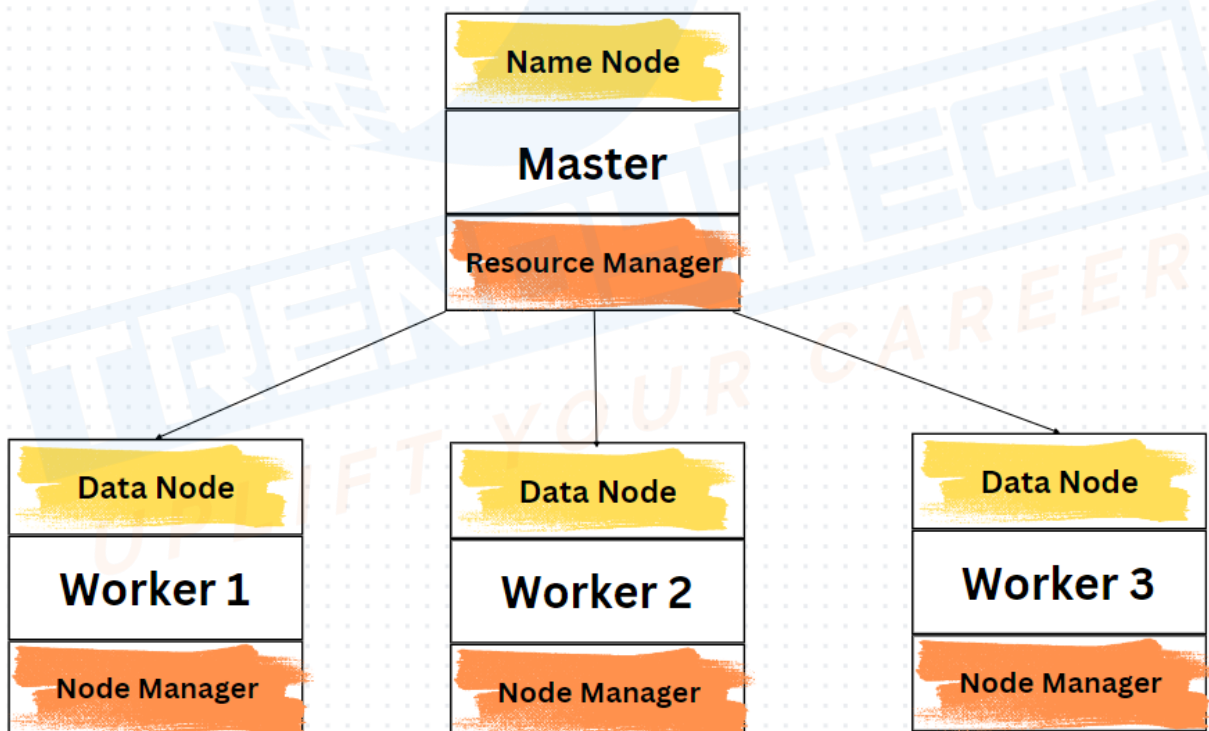
Hadoop Core Components :



YARN Architecture

YARN consists of two major components

1. Resource Manager (Master)
2. Node Manager (Worker / Slave)



**Name Node & Data Node
related to HDFS (STORAGE)**

**Resource Manager & Node Manager
related to YARN**

Processes involved in invoking a Hadoop Job from Client machine

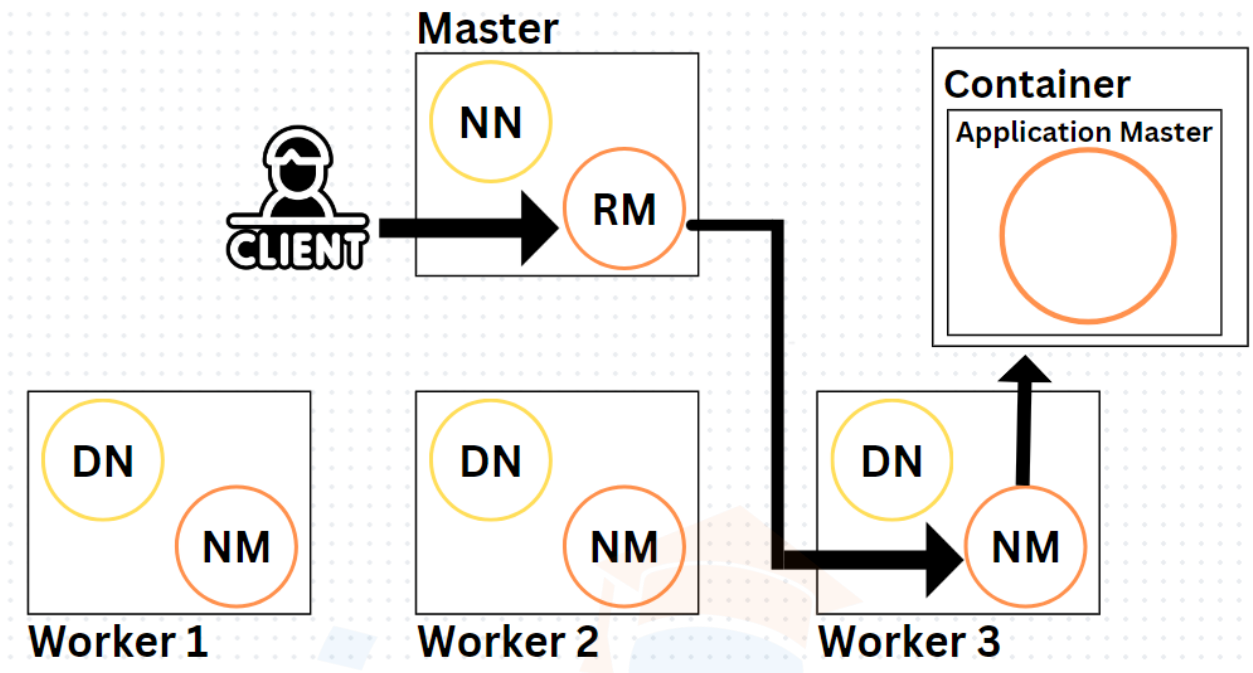
hadoop jar <Jar-Name>

On executing the above command on the client machine, following things happen :

1. **Request goes to Resource Manager**
2. **Resource Manager creates a Container in one of the Node Managers** (by coordinating with the Node Manager)
3. **An Application Master Service will be started inside this Container.**
(Application Master is a local Manager that manages the application)

Note : Application Master is responsible to negotiate for the required resources from the Resource Manager. It will interact with the Name Node to know where the Data is located and accordingly request for resources on specific nodes as it works on the principle of Data Locality.

Every Application has an Application Master i.e., If there are 20 Applications, then there would be 20 Application Masters running.



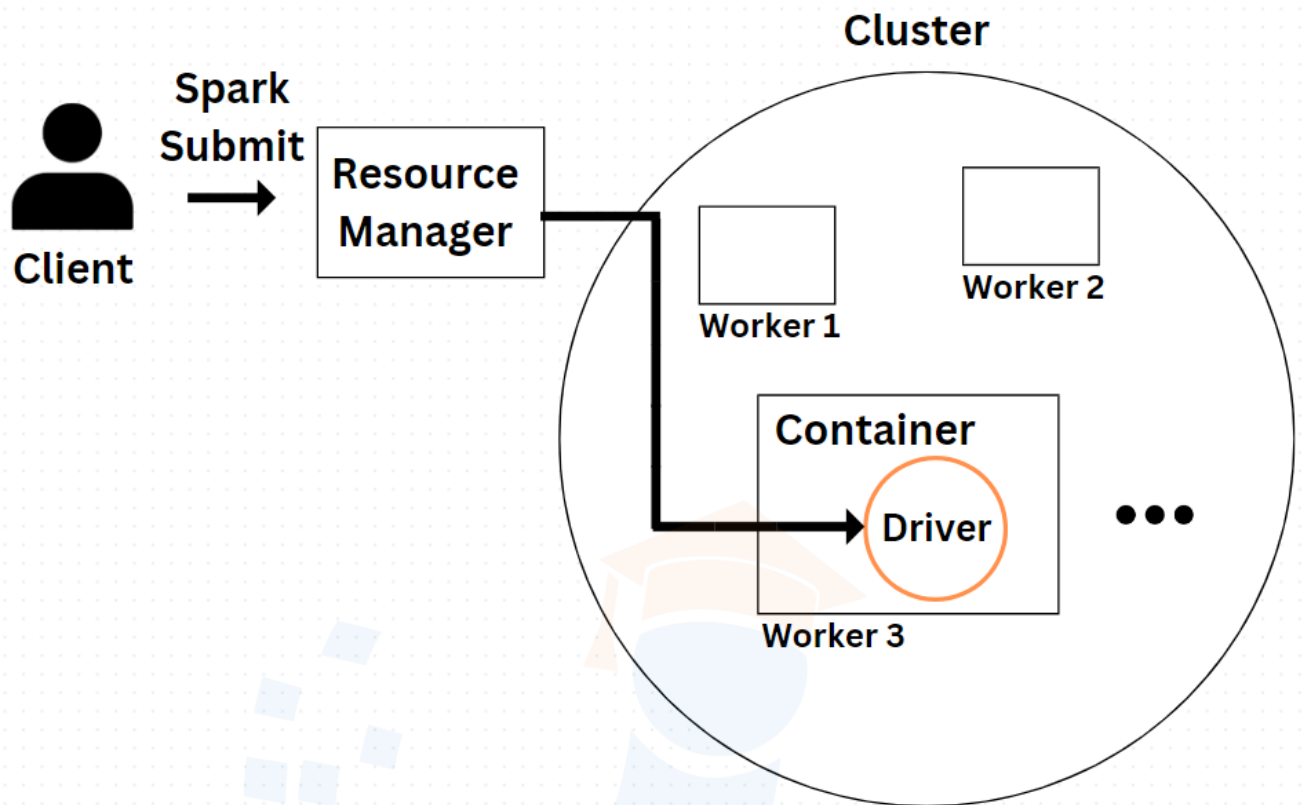
Uber Mode

Is a mode in which the application will be executed in the container where the application master is running. This mode is for scenarios where the job is small enough to be executed on a single container.

Two ways of running Spark Code-

1. Interactive Mode - NoteBooks / PySpark Shell
2. Bundle the code into a Jar and use Spark Submit to run the spark job.

Note : Every Spark Job has one Driver, Application Master acts like a driver which gets registered with the Resource Manager. If the Driver goes down, then the application crashes.



Two Modes in which Spark Runs-

1. **Client Mode** - It is an interactive mode where the intent is to view the results instantly. Notebooks / PySpark Shell are used for interactive mode.

Driver runs on the Client machine / Gateway node.

2. **Cluster Mode** - Code is packaged and submitted to the cluster for execution using Spark Submit.

Driver runs on the Cluster. Recommended for Production environment.

Resource Manager UI

Provides detailed information of the applications running and resources used like Application status, Memory, VCores...

hadoop

Cluster

About

Nodes

Node Labels

Applications

NEW

NEW SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used
58930	0	15	58915	58	101 GB

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes
3	0	1

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>

Show 20 entries

ID	User	Name	Application Type	Application Tags	Queue	Application Priority	StartTime	LaunchTime	FinishTime	Status
application_167599795986_59498	itv006879	pyspark-shell	SPARK		default	0	Mon Jun 12 20:06:34 +0550 2023	Mon Jun 12 20:06:35 +0550 2023	N/A	RUNNING

All Applications

Logged in as: dr.who

Running	Memory Used	Memory Total	Memory Reserved	Vcores Used	Vcores Total	Vcores Reserved
	101 GB	151.00 GB	0 B	58	90	0

Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
1	0	0	0	0

Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
es:1>	<memory:8192, vCores:4>	0

Search:

LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU Vcores	Allocated Memory MB	Reserved CPU Vcores	Reserved Memory MB	% of Queue	% of Cluster	Progress	Tracking UI	Blacklisted Nodes
Mon Jun 12 20:06:35 +0550 2023	N/A	RUNNING	UNDEFINED	3	3	5120	0	0	3.3	3.3	<div></div>	ApplicationMaster	0

Spark UI

spark 3.0.1

Jobs

Stages

Storage

Environment

Executors

SQL

pyspark-shell application UI

Spark Jobs (?)

User: itv006879

Total Uptime: 33 min

Scheduling Mode: FIFO

Completed Jobs: 3

Event Timeline

Enable zooming

Executors

Added

Removed

Jobs

Succeeded

Failed

Running

Mon 12 June

10:36

10:37

10:38

10:39

10:40

10:41

10:42

10:43

10:44

10:45

10:46

Completed Jobs (3)

Page: 1

1 Pages. Jump to 1. Show 100 items in a page. Go

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
2	showString at NativeMethodAccessorImpl.java:0 showString at NativeMethodAccessorImpl.java:0	2023/06/12 10:46:07	26 ms	1/1	1/1

Accessing Columns in PySpark

1. String Notation

Ex : `df.select("*").show()`

2. Prefixing Column name with Dataframe name

Ex : `df.order_date`

3. Array Notation

Ex : `df['order_date']`

4. Column Object Notation

Ex : `column('order_date')` or `col('order_date')`

5. Column Expression

Ex : `expr("order_date")`

Why are there so many ways of accessing columns?

- Prefixing Column name with Dataframe name

If two different dataframes have columns with the same name, then this would lead to ambiguity as the system is not aware of which dataframe column to choose.

Ex - Say `cust_id` is the column available in both orders and customer dataframes. Then the system is not sure whether to pick `cust_id` of customer dataframe or orders dataframe. Therefore, prefix the dataframe name with the column name to prevent any ambiguity.

- Column Expression

Is required when evaluation needs to be performed in a SQL way.

Ex - Say you need to increment customer ID and create a new customer ID

`expr("cust_id + 1 as new_cust_id")`

- Column Object

Provides various predefined functions to achieve the desired results in a programmatic approach.

Ex -

```
orders_df.select("*").where(col('order_status').like('PENDING%')).show()
```

OR

```
orders_df.select("*").where("order_status like 'PENDING%' ").show()
```

Aggregate Functions

Combining multiple input rows together to give a consolidated output.

Simple Aggregations Example

Consider you have a orders.csv dataset and you are required to

- Count the total number of records
- Count number of distinct invoice ids
- Sum of Quantities
- Average unit price

Lets see 3 ways of solving the above

1. Programmatic style

Create and load a dataframe

```
from pyspark.sql.functions import *  
  
orders_df.select(count("*").alias("row_count"),  
countDistinct("invoiceno").alias("unique_invoice"),  
sum("quantity").alias("total_quantity"),  
avg("unitprice").alias("avg_price")).show()
```


2. Column Expression style

```
orders_df.selectExpr("count(*) as row_count",
"count(distinct(invoiceno)) as unique_invoice", "sum(quantity) as
total_quantity", "avg(unitprice) as avg_price").show()
```

3. Spark SQL style

Create a Orders table

```
spark.sql("select count(*) as row_count, count(distinct(invoiceno))
as unique_invoice, sum(quantity) as total_quantity, avg(unitprice)
as avg_price from orders").show()
```

Grouping Aggregations Example

Consider you have a orders.csv dataset and you are required to group based on invoice number and country.

- Find the total quantity for each group
- Find the total invoice amount (Amount = Quantity * UnitPrice)

Lets see 3 ways of solving the above

1. Programmatic style

Create and load dataframe with the given dataset

```
from pyspark.sql.functions import *

summary_df = orders_df \
.groupBy("country", "invoiceno") \
.agg(sum("quantity").alias("total_quantity"), sum(expr("quantity *
unitprice")).alias("invoice_value")).sort("invoiceno")
```


2. Column Expression style

```
summary_df = orders_df \
    .groupBy("country","invoiceno") \
    .agg(expr("sum(quantity) as total_quantity"), expr("sum(quantity *
    unitprice) as invoice_value")).sort("invoiceno")
```

3. Spark SQL style

Create a Orders table

```
orders_df.createOrReplaceTempView("orders")

spark.sql(""" select country, invoiceno, sum(quantity) as
total_quantity, sum(quantity * unitprice) as invoice_value from
orders group by country, invoiceno order by invoiceno """).show()
```

Windowing Aggregations Example

Consider you have a windowdata.csv under /public/trendytech/datasets/windowdata.csv in the lab dataset and you are required to define the following 3 parameters -

1. **Partition Column** Partition by based on country
2. **Sorting Column** Sort based on week number
3. **Window Size** (Define the size by mentioning the start row and end row)

- Find the running total of invoice value

Create and load dataframe with the given dataset

```
from pyspark.sql.functions import *
```

```
mywindow = Window.partitionBy("country") \
    .orderBy("weeknum") \
```

```
.rowsBetween(Window.unboundedPreceding, Window.currentRow)
```

```
result_df =
```

```
orders_df.withColumn("running_total",sum("invoicevalue"))
```

```
.over(mywindow))
```

```
result_df.show()
```

Simple Aggregations	Grouping Aggregations	Windowing Aggregations
Accepts multiple input rows to give only one output row	Accepts multiple input rows belonging to a group (grouping is done on a column). For each group, there is one output row	Output generated by performing operations on a predefined rows set within a window
SUM , COUNT, DISTINCT	groupBy	partitionBy, orderBy, rowsBetween

Windowing Functions

1. rank
2. dense_rank
3. row_number
4. lead
5. lag

Consider you have a windowdatamodified.csv under **/public/trendytech/datasets/windowdatamodified.csv** in the lab.

This csv file has some values repeating to demonstrate the behaviour of the above windowing functions.

Rank

```
from pyspark.sql import SparkSession
import getpass
username = getpass.getuser()
spark = SparkSession. \
    builder. \
    config('spark.ui.port', '0'). \
    config("spark.sql.warehouse.dir", f"/user/{username}/warehouse"). \
    enableHiveSupport(). \
    master('yarn'). \
    getOrCreate()
```

```
orders_df = spark.read \
    .format("csv") \
    .option("inferSchema", "true") \
    .option("header", "true") \
    .load("/public/trendytech/datasets/windowdatamodified.csv")
```

```
from pyspark.sql import *
from pyspark.sql.functions import *
```

```
mywindow = Window.partitionBy("country") \
    .orderBy(desc("invoicevalue"))
```

```
results_df = orders_df.withColumn("rank", rank().over(mywindow))
```

```
results_df.show()
```

country	weeknum	numinvoices	totalquantity	invoicevalue	rank
Sweden	50	3	3714	2646.3	1
Germany	49	12	1852	1800.0	1
Germany	50	15	1973	1800.0	1
Germany	48	11	1795	1600.0	3
Germany	51	5	1103	1600.0	3
France	50	6	529	537.32	1
France	51	5	847	500.0	2
France	49	9	2303	500.0	2
France	48	4	1299	500.0	2
Belgium	48	1	528	800.0	1
Belgium	51	2	942	800.0	1
Belgium	50	2	285	625.16	3
Finland	50	1	1254	892.8	1
India	49	5	1280	3284.1	1
India	50	5	1184	2321.78	2
India	51	5	95	300.0	3
India	48	7	2822	300.0	3
Italy	48	1	164	427.8	1
Italy	51	1	131	383.7	2
Italy	49	1	-2	-17.0	3

only showing top 20 rows

Dense Rank

```
results_df = orders_df.withColumn("rank",dense_rank().over(mywindow))
```

```
results_df.show()
```

```
+-----+-----+-----+-----+-----+-----+
|country|weeknum|numinvoices|totalquantity|invoicevalue|rank|
+-----+-----+-----+-----+-----+-----+
|Sweden|50|3|3714|2646.3|1|
|Germany|49|12|1852|1800.0|1|
|Germany|50|15|1973|1800.0|1|
|Germany|48|11|1795|1600.0|2|
|Germany|51|5|1103|1600.0|2|
|France|50|6|529|537.32|1|
|France|51|5|847|500.0|2|
|France|49|9|2303|500.0|2|
|France|48|4|1299|500.0|2|
|Belgium|48|1|528|800.0|1|
|Belgium|51|2|942|800.0|1|
|Belgium|50|2|285|625.16|2|
|Finland|50|1|1254|892.8|1|
|India|49|5|1280|3284.1|1|
|India|50|5|1184|2321.78|2|
|India|51|5|95|300.0|3|
|India|48|7|2822|300.0|3|
|Italy|48|1|164|427.8|1|
|Italy|51|1|131|383.7|2|
|Italy|49|1|-2|-17.0|3|
+-----+-----+-----+-----+-----+-----+
```

only showing top 20 rows

Note :

In Rank, some ranks can be skipped if there are clashes in the ranks.

In Dense Rank, the ranks are not skipped even if there are clashes in the ranks.

In Row Number, different row numbers are assigned even in case of a tie. It plays an important role in calculating the top-n results.

Consider an Example Scenario of how ranks are assigned to the students based on the marks scored

Student Name	Marks Scored	rank	dense-rank	row-number
Ankur	100	1	1	1
Satish	100	1	1	2
Kapil	100	1	1	3
Kaushik	99	4	2	4
Ram	99	4	2	5
Rohit	98	6	3	6

Row Number

```
results_df = orders_df.withColumn("rank",row_number().over(mywindow))
```

```
results_df.show()
```

```
+-----+-----+-----+-----+-----+-----+
|country|weeknum|numinvoices|totalquantity|invoicevalue|rank|
+-----+-----+-----+-----+-----+-----+
|Sweden|50|3|3714|2646.3|1|
|Germany|49|12|1852|1800.0|1|
|Germany|50|15|1973|1800.0|2|
|Germany|48|11|1795|1600.0|3|
|Germany|51|5|1103|1600.0|4|
|France|50|6|529|537.32|1|
|France|51|5|847|500.0|2|
|France|49|9|2303|500.0|3|
|France|48|4|1299|500.0|4|
|Belgium|48|1|528|800.0|1|
|Belgium|51|2|942|800.0|2|
|Belgium|50|2|285|625.16|3|
|Finland|50|1|1254|892.8|1|
|India|49|5|1280|3284.1|1|
|India|50|5|1184|2321.78|2|
|India|51|5|95|300.0|3|
|India|48|7|2822|300.0|4|
|Italy|48|1|164|427.8|1|
|Italy|51|1|131|383.7|2|
|Italy|49|1|-2|-17.0|3|
+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

Note : When you need to compare two rows, then the lead or lag function is used.

Lead - Is used when the current row needs to be compared with the next row.

Lag - Is used when the current row needs to be compared with the previous row.

```
from pyspark.sql import SparkSession
import getpass
username = getpass.getuser()
spark= SparkSession. \
builder. \
config('spark.ui.port','0'). \
config("spark.sql.warehouse.dir", f"/user/{username}/warehouse"). \
enableHiveSupport(). \
master('yarn'). \
getOrCreate()
```

```
orders_df = spark.read \
.format("csv") \
.option("inferSchema","true") \
.option("header","true") \
.load("/public/trendytech/datasets/windowdatamodified.csv")
```

```
from pyspark.sql import *
from pyspark.sql.functions import *
```

```
mywindow = Window.partitionBy("country") \
.orderBy("weeknum")
```

```
results_df = orders_df.withColumn("previous_week",lag("invoicevalue").over(mywindow))
```

```
results_df.show()
```

country	weeknum	numinvoiced	totalquantity	invoicevalue	previous_week
Sweden	50	3	3714	2646.3	null
Germany	48	11	1795	1600.0	null
Germany	49	12	1852	1800.0	1600.0
Germany	50	15	1973	1800.0	1800.0
Germany	51	5	1103	1600.0	1800.0
France	48	4	1299	500.0	null
France	49	9	2303	500.0	500.0
France	50	6	529	537.32	500.0
France	51	5	847	500.0	537.32
Belgium	48	1	528	800.0	null
Belgium	50	2	285	625.16	800.0
Belgium	51	2	942	800.0	625.16
Finland	50	1	1254	892.8	null
India	48	7	2822	300.0	null
India	49	5	1280	3284.1	300.0
India	50	5	1184	2321.78	3284.1
India	51	5	95	300.0	2321.78
Italy	48	1	164	427.8	null
Italy	49	1	-2	-17.0	427.8
Italy	51	1	131	383.7	-17.0

only showing top 20 rows

Analysing log files to find some valuable Inferences.

1. First develop the logic with sample data and then apply to original data.

`//Create a Spark Session`

```
logs_data = [("INFO","2015-8-8 20:49:22"),
("WARN","2015-1-14 20:05:00"),
("INFO","2017-6-14 00:08:35"),
("INFO","2016-1-18 11:50:14"),
("DEBUG","2017-7-1 12:55:02"),
("INFO","2014-2-26 12:34:21"),
("INFO","2015-7-12 11:13:47"),
("INFO","2017-4-15 01:20:18"),
("DEBUG","2016-11-2 20:19:23"),
("INFO","2012-8-20 10:09:44")] //Sample log Data
```

```
log_df = spark.createDataFrame(logs_data).toDF('log_level','log_time')
//creating dataframe from local hard-coded data
```

```
new_log_df = log_df.withColumn("logtime", to_timestamp("log_time"))
//changing the datatype of the column log_time from string to timestamp
```

```
//In order to operate on the data like a sql table create a temp view table.
new_log_df.createOrReplaceTempView("serverlogs")
```

```
spark.sql(select loglevel, date_format(logtime, 'MMMM') as month from
serverlogs").show()
```

`//if only the month needs to be extracted.`

```
spark.sql(select loglevel, date_format(logtime, 'MMMM') as month,
count(*) as total_occurrence from serverlogs group by loglevel,
month").show() //Applying transformations - total occurrence of different log
status like WARN,INFO ... in the log file by grouping based on loglevel and
month
```


2. Now working on the original dataset after ensuring that the logic is functional without any errors.

Plug in the main dataset

```
logschema = "loglevel string, logtime timestamp"
```

```
log_df = spark.read \
    .format("csv") \
    .schema(logschema) \
    .load("/public/trendytech/datasets/logsdata1m.csv")
```

//Now the same aggregations as in the previous case can be applied to this dataframe.

Pivot Table

Provides a more intuitive view by which data can be easily analysed for insights

```
spark.sql("select loglevel, date_format(logtime, 'MMMM') as month from serverlogs").groupBy('loglevel').pivot('month').count.show()
```

Pivot Table View

	Month ->							
	January	February	March	April	May	June	● ● ●	December
Error	1000	1200						
Info	100	89						
Warn	100	99						
Fatal	100	100						
Debug	99	79						

Note :

Optimization - By explicitly providing the list of values on the pivot column, the system will not be scanning the data to create the list of pivot values. This will save some processing time and improve the performance of query execution.

```
month_list = ['Jan', 'Feb', 'Mar', ... 'dec']
```

```
spark.sql("select loglevel, date_format(logtime, 'MMMM') as month from serverlogs").groupBy('loglevel').pivot('month', month_list).count.show()
```